



UNIVERSITY ECHAHD HAMA LAKHDAR, EL-OUED
INSTITUTE OF EXACT SCIENCES
DEPARTMENT OF COMPUTER SCIENCE
2^{ed} Master :
ARTIFICIAL INTELLIGENCE AND DATA SCIENCE
Semester : 3.2025

Lab 6 : Object Detection using R-CNN From Region Proposals to Deep Feature Classification Master Level – Computer Vision

Introduction

This lab session aims to teach step-by-step implementation of the R-CNN (Region-based Convolutional Neural Network) architecture using Python. You will explore the process of generating region proposals, extracting CNN features, classifying them, and evaluating detections. A subset of the COCO dataset will be used as a case study, with an optional extension for vignette (drug sticker) data.

1 Objective

- Understand R-CNN architecture components.
- Implement Selective Search for region proposals.
- Extract CNN features using pre-trained networks.
- Train and evaluate SVM classifiers.
- Apply Non-Maximum Suppression (NMS).
- Evaluate performance on COCO and vignette datasets.

2 Task 1 : Dataset Preparation

COCO Subset

Download a subset of COCO (e.g., categories : person, dog, car). Resize images to 416×416 for uniformity.

Listing 1 – Load and prepare COCO subset

```
1 from pycocotools.coco import COCO
2 import requests, os, cv2
3
4 dataDir = 'coco'
5 dataType = 'val2017'
6 annFile = f'{dataDir}/annotations/instances_{dataType}.json'
7
8 coco = COCO(annFile)
9 catIds = coco.getCatIds(catNms=['person','dog','car'])
10 imgIds = coco.getImgIds(catIds=catIds)
11
12 img = coco.loadImgs(imgIds[0])[0]
13 img_data = requests.get(img['coco_url']).content
14 with open('sample.jpg', 'wb') as f:
15     f.write(img_data)
```

Questions :

-
- a) Why is resizing necessary for CNN feature extraction ?
 - b) What challenges arise from using only a subset of COCO ?

Vignette Dataset (Extension Questions)

- a) How does lighting variation in vignette images affect region proposals ?
 - b) Which preprocessing methods could improve OCR detection ?
-

3 Task 2 : Region Proposals using Selective Search

Listing 2 – Generate region proposals with Selective Search

```

1 import selectivesearch
2 import matplotlib.pyplot as plt
3
4 img = cv2.imread('sample.jpg')
5 img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
6 _, regions = selectivesearch.selective_search(img_rgb, scale=500, sigma=0.9,
7 min_size=50)
8
9 candidates = []
10 for r in regions:
11     if r['rect'] in candidates:
12         continue
13     x, y, w, h = r['rect']
14     if w * h < 2000:
15         continue
16     candidates.append(r['rect'])
17
18 for (x, y, w, h) in candidates[:50]:
19     cv2.rectangle(img_rgb, (x, y), (x+w, y+h), (255, 0, 0), 2)
20
21 plt.imshow(img_rgb)
22 plt.title('Region Proposals')
23 plt.show()

```

Questions :

- a) What are the trade-offs in adjusting `scale`, `sigma`, and `min_size` ?
- b) Why does Selective Search tend to over-generate proposals ?

Vignette Dataset (Extension Questions)

- a) How does Selective Search handle text-dense vignette images ?
 - b) Could MSER yield better proposals for text ? Why ?
-

4 Task 3 : CNN Feature Extraction

Use a pre-trained CNN (e.g., VGG16) to extract deep features from proposed regions.

Listing 3 – Extract CNN features using VGG16

```

1 from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input
2 from tensorflow.keras.preprocessing.image import img_to_array
3 import numpy as np
4
5 model = VGG16(weights='imagenet', include_top=False, pooling='avg')

```

```

6 features = []
7 for (x, y, w, h) in candidates[:200]:
8     roi = img_rgb[y:y+h, x:x+w]
9     roi = cv2.resize(roi, (224, 224))
10    roi = img_to_array(roi)
11    roi = np.expand_dims(roi, axis=0)
12    roi = preprocess_input(roi)
13    f = model.predict(roi)
14    features.append(f.flatten())
15

```

Questions :

- a) Which CNN layers provide the most robust object features?
- b) How does feature dimensionality affect SVM training time?

Vignette Dataset (Extension Questions)

- a) Which CNN features best capture character shapes?
 - b) Would fine-tuning on text patches improve OCR accuracy?
-

5 Task 4 : SVM Classification

Train an SVM classifier on the extracted features.

Listing 4 – Train an SVM classifier

```

1 from sklearn.svm import LinearSVC
2 from sklearn.model_selection import train_test_split
3
4 X = np.array(features)
5 y = np.random.randint(0, 2, len(X)) # Placeholder labels
6
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
8 svm = LinearSVC(max_iter=10000)
9 svm.fit(X_train, y_train)
10
11 print("Accuracy:", svm.score(X_test, y_test))

```

Questions :

- a) Why is an SVM used instead of softmax here?
- b) How could we handle class imbalance in object detection?

Vignette Dataset (Extension Questions)

- a) How would you label “entête” regions versus background?
 - b) Can SVMs separate characters from background efficiently?
-

6 Task 5 : Detection and Non-Maximum Suppression

Apply SVM predictions to region proposals and refine detections with NMS.

Listing 5 – Apply NMS to reduce overlapping detections

```

1 from imutils.object_detection import non_max_suppression
2
3 boxes = np.array([[x, y, x+w, y+h] for (x,y,w,h) in candidates[:200]])
4 scores = np.random.rand(len(boxes)) # Simulated scores

```

```
5 final_boxes = non_max_suppression(boxes, probs=scores, overlapThresh=0.3)
6 print("Detections after NMS:", len(final_boxes))
7
```

Questions :

- a) What effect does the IoU threshold have on detections ?
- b) How could overlapping boxes cause false positives ?

Vignette Dataset (Extension Questions)

- a) How can NMS be tuned for dense character regions ?
- b) What IoU threshold suits OCR-style text boxes ?

7 Task 6 : Evaluation and Visualization

Evaluate the system using mean Average Precision (mAP).

Listing 6 – Evaluate detection performance

```
1 from sklearn.metrics import average_precision_score
2
3 y_true = np.random.randint(0, 2, len(scores))
4 mAP = average_precision_score(y_true, scores)
5 print("Mean Average Precision (mAP):", mAP)
```

Questions :

- a) Why is mAP a better metric than accuracy for detection ?
- b) How do localization errors impact mAP ?

Vignette Dataset (Extension Questions)

- a) Which metric suits character-level accuracy best ?
- b) How can we evaluate both detection and OCR jointly ?

8 Task 7 : Backbone Variation (Optional)

Try different CNN backbones for feature extraction (ResNet, MobileNet).

Listing 7 – Use ResNet backbone for feature extraction

```
1 from tensorflow.keras.applications import ResNet50
2
3 resnet = ResNet50(weights='imagenet', include_top=False, pooling='avg')
4 features_resnet = [resnet.predict(preprocess_input(np.expand_dims(cv2.resize(
    img_rgb[y:y+h, x:x+w], (224,224)), axis=0))).flatten() for (x,y,w,h) in
    candidates[:100]]
```

Questions :

- a) How does backbone depth affect computation time ?
- b) What trade-offs exist between MobileNet and ResNet ?

Vignette Dataset (Extension Questions)

- a) Which backbone best captures small-text details ?
- b) How could OCR-specific backbones improve detection ?

9 Conclusion

Students have implemented a complete R-CNN pipeline step by step :

1. Region Proposals (Selective Search)
2. CNN Feature Extraction
3. SVM Classification
4. Non-Maximum Suppression
5. Evaluation

By extending the same workflow to vignette data, students can adapt R-CNN principles for OCR and printed-text detection.